

**NAME**

perfmonctl – interface to PMU

**SYNOPSIS**

```
#include <syscall.h>
#include <perfmon.h>
```

```
long perfmonctl(int fd, int cmd, void *arg, int nargs);
```

**DESCRIPTION**

**perfmonctl** system call provides an interface to PMU (performance monitoring unit). PMU consists of PMD (performance monitoring data) registers and PMC (performance monitoring control) registers, where are gathered the hardware statistic.

**perfmonctl** will apply a function *cmd* to input arguments *arg*. The number of arguments is defined by input variable *narg*. *fd* specifies the perfmon context to operate on.

The implemented commands *cmd* are:

**PFM\_CREATE\_CONTEXT**

set up a context

```
perfmonctl(int fd, PFM_CREATE_CONTEXT, pfarg_context_t *ctxt, 1);
```

The *fd* parameter is ignored. A new context is created as specified in *ctxt* and its file descriptor is returned in *ctxt->ctx\_fd*.

The file descriptor, apart from passing it to **perfmonctl**, can be used to read event notifications (type **pfm\_msg\_t**) using the **read(2)** system call. Both **select(2)** and **poll(2)** can be used to wait for event notifications.

The context can be destroyed using the **close(2)** system call.

**PFM\_WRITE\_PMCS**

set PMC registers

```
perfmonctl(int fd, PFM_WRITE_PMCS, pfarg_pmc_t *pmcs, n);
```

**PFM\_WRITE\_PMDs**

set PMD registers

```
perfmonctl(int fd, PFM_WRITE_PMDs, pfarg_pmd_t *pmds, n);
```

**PFM\_READ\_PMDs**

read PMD registers

```
perfmonctl(int fd, PFM_READ_PMDs, pfarg_pmd_t *pmds, n);
```

**PFM\_START**

start monitoring

```
perfmonctl(int fd, PFM_START, arg, 1);
perfmonctl(int fd, PFM_START, NULL, 0);
```

**PFM\_STOP**

stop monitoring

```
perfmonctl(int fd, PFM_STOP, NULL, 0);
```

**PFM\_LOAD\_CONTEXT**

attach the context to a thread

```
perfmonctl(int fd, PFM_LOAD_CONTEXT, pfarg_load_t *largs, 1);
```

**PFM\_UNLOAD\_CONTEXT**

detach the context from a thread

```
perfmonctl(int fd, PFM_UNLOAD_CONTEXT , NULL , 0);
```

**PFM\_RESTART**

restart monitoring after receiving an overflow notification

```
perfmonctl(int fd, PFM_RESTART , NULL , 0);
```

**PFM\_CREATE\_EVTSETS**

create or modify event sets

```
perfmonctl(int fd, PFM_CREATE_EVTSETS, pfarg_setdesc_t *desc , n);
```

**PFM\_DELETE\_EVTSETS**

delete event sets

```
perfmonctl(int fd, PFM_DELETE_EVTSET, pfarg_setdesc_t *desc , n);
```

**PFM\_GETINFO\_EVTSETS**

get information about event sets

```
perfmonctl(int fd, PFM_GETINFO_EVTSETS, pfarg_setinfo_t *info, n);
```

**RETURN VALUE**

**performctl** returns zero when the operation is successful. On error -1 is returned and an error code is set in **errno**.

**AVAILABILITY**

This syscall is implemented only on the IA-64 architecture since kernel 2.6.

**SEE ALSO**

**gprof(1)**, The perfmon2 interface specification