

NAME

`recv`, `recvfrom`, `recvmsg` – receive a message from a socket

SYNOPSIS

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
ssize_t recv(int sockfd, void *buf, size_t len, int flags);
```

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags,
                 struct sockaddr *src_addr, socklen_t *addrlen);
```

```
ssize_t recvmsg(int sockfd, struct msghdr *msg, int flags);
```

DESCRIPTION

The `recvfrom()` and `recvmsg()` calls are used to receive messages from a socket, and may be used to receive data on a socket whether or not it is connection-oriented.

If `src_addr` is not NULL, and the underlying protocol provides the source address, this source address is filled in. When `src_addr` is NULL, nothing is filled in; in this case, `addrlen` is not used, and should also be NULL. The argument `addrlen` is a value-result argument, which the caller should initialize before the call to the size of the buffer associated with `src_addr`, and modified on return to indicate the actual size of the source address. The returned address is truncated if the buffer provided is too small; in this case, `addrlen` will return a value greater than was supplied to the call.

The `recv()` call is normally used only on a *connected* socket (see `connect(2)`) and is identical to `recvfrom()` with a NULL `src_addr` argument.

All three routines return the length of the message on successful completion. If a message is too long to fit in the supplied buffer, excess bytes may be discarded depending on the type of socket the message is received from.

If no messages are available at the socket, the receive calls wait for a message to arrive, unless the socket is non-blocking (see `fcntl(2)`), in which case the value `-1` is returned and the external variable `errno` is set to **EAGAIN** or **EWOULDBLOCK**. The receive calls normally return any data available, up to the requested amount, rather than waiting for receipt of the full amount requested.

The `select(2)` or `poll(2)` call may be used to determine when more data arrives.

The `flags` argument to a `recv()` call is formed by *OR*'ing one or more of the following values:

MSG_CMSG_CLOEXEC (`recvmsg()` only; since Linux 2.6.23)

Set the close-on-exec flag for the file descriptor received via a Unix domain file descriptor using the **SCM_RIGHTS** operation (described in `unix(7)`). This flag is useful for the same reasons as the **O_CLOEXEC** flag of `open(2)`.

MSG_DONTWAIT (since Linux 2.2)

Enables non-blocking operation; if the operation would block, the call fails with the error **EAGAIN** or **EWOULDBLOCK** (this can also be enabled using the **O_NONBLOCK** flag with the **F_SETFL** `fcntl(2)`).

MSG_ERRQUEUE (since Linux 2.2)

This flag specifies that queued errors should be received from the socket error queue. The error is passed in an ancillary message with a type dependent on the protocol (for IPv4 **IP_RECVERR**). The user should supply a buffer of sufficient size. See `cmsg(3)` and `ip(7)` for more information. The payload of the original packet that caused the error is passed as normal data via `msg_iovec`. The original destination address of the datagram that caused the error is supplied via `msg_name`.

For local errors, no address is passed (this can be checked with the `cmsg_len` member of the `cmsghdr`). For error receives, the **MSG_ERRQUEUE** is set in the `msghdr`. After an error has been passed, the pending socket error is regenerated based on the next queued error and will be passed on the next socket operation.

The error is supplied in a *sock_extended_err* structure:

```
#define SO_EE_ORIGIN_NONE  0
#define SO_EE_ORIGIN_LOCAL 1
#define SO_EE_ORIGIN_ICMP  2
#define SO_EE_ORIGIN_ICMP6 3

struct sock_extended_err
{
    uint32_t ee_errno; /* error number */
    uint8_t ee_origin; /* where the error originated */
    uint8_t ee_type; /* type */
    uint8_t ee_code; /* code */
    uint8_t ee_pad; /* padding */
    uint32_t ee_info; /* additional information */
    uint32_t ee_data; /* other data */
    /* More data may follow */
};

struct sockaddr *SO_EE_OFFENDER(struct sock_extended_err *);
```

ee_errno contains the *errno* number of the queued error. *ee_origin* is the origin code of where the error originated. The other fields are protocol-specific. The macro **SOCK_EE_OFFENDER** returns a pointer to the address of the network object where the error originated from given a pointer to the ancillary message. If this address is not known, the *sa_family* member of the *sockaddr* contains **AF_UNSPEC** and the other fields of the *sockaddr* are undefined. The payload of the packet that caused the error is passed as normal data.

For local errors, no address is passed (this can be checked with the *cmsg_len* member of the *cmsg_hdr*). For error receives, the **MSG_ERRQUEUE** is set in the *msghdr*. After an error has been passed, the pending socket error is regenerated based on the next queued error and will be passed on the next socket operation.

MSG_OOB

This flag requests receipt of out-of-band data that would not be received in the normal data stream. Some protocols place expedited data at the head of the normal data queue, and thus this flag cannot be used with such protocols.

MSG_PEEK

This flag causes the receive operation to return data from the beginning of the receive queue without removing that data from the queue. Thus, a subsequent receive call will return the same data.

MSG_TRUNC (since Linux 2.2)

For raw (**AF_PACKET**), Internet datagram (since Linux 2.4.27/2.6.8), and netlink (since Linux 2.6.22) sockets: return the real length of the packet or datagram, even when it was longer than the passed buffer. Not implemented for Unix domain sockets.

For use with Internet stream sockets, see **tcp(7)**.

MSG_WAITALL (since Linux 2.2)

This flag requests that the operation block until the full request is satisfied. However, the call may still return less data than requested if a signal is caught, an error or disconnect occurs, or the next data to be received is of a different type than that returned.

The **recvmsg()** call uses a *msghdr* structure to minimize the number of directly supplied arguments. This structure is defined as follows in *<sys/socket.h>*:

```
struct iovec {
    void *iov_base; /* Scatter/gather array items */
    /* Starting address */
```

```

    size_t iov_len;          /* Number of bytes to transfer */
};

struct msghdr {
    void      *msg_name;     /* optional address */
    socklen_t  msg_namelen;  /* size of address */
    struct iovec *msg_iov;    /* scatter/gather array */
    size_t     msg_iovlen;   /* # elements in msg_iov */
    void      *msg_control;   /* ancillary data, see below */
    socklen_t  msg_controllen; /* ancillary data buffer len */
    int        msg_flags;     /* flags on received message */
};

```

Here *msg_name* and *msg_namelen* specify the source address if the socket is unconnected; *msg_name* may be given as a null pointer if no names are desired or required. The fields *msg_iov* and *msg_iovlen* describe scatter-gather locations, as discussed in **readv**(2). The field *msg_control*, which has length *msg_controllen*, points to a buffer for other protocol control-related messages or miscellaneous ancillary data. When **recvmsg**() is called, *msg_controllen* should contain the length of the available buffer in *msg_control*; upon return from a successful call it will contain the length of the control message sequence.

The messages are of the form:

```

struct cmsghdr {
    socklen_t  cmsg_len;     /* data byte count, including hdr */
    int        cmsg_level;   /* originating protocol */
    int        cmsg_type;    /* protocol-specific type */
    /* followed by
       unsigned char cmsg_data[]; */
};

```

Ancillary data should only be accessed by the macros defined in **cmsg**(3).

As an example, Linux uses this auxiliary data mechanism to pass extended errors, IP options or file descriptors over Unix sockets.

The *msg_flags* field in the *msghdr* is set on return of **recvmsg**(). It can contain several flags:

MSG_EOR

indicates end-of-record; the data returned completed a record (generally used with sockets of type **SOCK_SEQPACKET**).

MSG_TRUNC

indicates that the trailing portion of a datagram was discarded because the datagram was larger than the buffer supplied.

MSG_CTRUNC

indicates that some control data were discarded due to lack of space in the buffer for ancillary data.

MSG_OOB

is returned to indicate that expedited or out-of-band data were received.

MSG_ERRQUEUE

indicates that no data was received but an extended error from the socket error queue.

RETURN VALUE

These calls return the number of bytes received, or -1 if an error occurred. The return value will be 0 when the peer has performed an orderly shutdown.

ERRORS

These are some standard errors generated by the socket layer. Additional errors may be generated and returned from the underlying protocol modules; see their manual pages.

EAGAIN or EWOULDBLOCK

The socket is marked non-blocking and the receive operation would block, or a receive timeout had been set and the timeout expired before data was received. POSIX.1-2001 allows either error to be returned for this case, and does not require these constants to have the same value, so a portable application should check for both possibilities.

EBADF

The argument *sockfd* is an invalid descriptor.

ECONNREFUSED

A remote host refused to allow the network connection (typically because it is not running the requested service).

EFAULT

The receive buffer pointer(s) point outside the process's address space.

EINTR

The receive was interrupted by delivery of a signal before any data were available; see **signal(7)**.

EINVAL

Invalid argument passed.

ENOMEM

Could not allocate memory for **recvmsg()**.

ENOTCONN

The socket is associated with a connection-oriented protocol and has not been connected (see **connect(2)** and **accept(2)**).

ENOTSOCK

The argument *sockfd* does not refer to a socket.

CONFORMING TO

4.4BSD (these function calls first appeared in 4.2BSD), POSIX.1-2001.

POSIX.1-2001 only describes the **MSG_OOB**, **MSG_PEEK**, and **MSG_WAITALL** flags.

NOTES

The prototypes given above follow glibc2. The Single Unix Specification agrees, except that it has return values of type *ssize_t* (while 4.x BSD and libc4 and libc5 all have *int*). The *flags* argument is *int* in 4.x BSD, but *unsigned int* in libc4 and libc5. The *len* argument is *int* in 4.x BSD, but *size_t* in libc4 and libc5. The *addrlen* argument is *int ** in 4.x BSD, libc4 and libc5. The present *socklen_t ** was invented by POSIX. See also **accept(2)**.

According to POSIX.1-2001, the *msg_controllen* field of the *msghdr* structure should be typed as *socklen_t*, but glibc currently (2.4) types it as *size_t*.

EXAMPLE

An example of the use of **recvfrom()** is shown in **getaddrinfo(3)**.

SEE ALSO

fcntl(2), **getsockopt(2)**, **read(2)**, **select(2)**, **shutdown(2)**, **socket(2)**, **cmsg(3)**, **socketatmark(3)**, **socket(7)**

COLOPHON

This page is part of release 3.22 of the Linux *man-pages* project. A description of the project, and information about reporting bugs, can be found at <http://www.kernel.org/doc/man-pages/>.