

NAME

splice – splice data to/from a pipe

SYNOPSIS

```
#define _GNU_SOURCE
#include <fcntl.h>
```

```
long splice(int fd_in, loff_t *off_in, int fd_out,
            loff_t *off_out, size_t len, unsigned int flags);
```

DESCRIPTION

splice() moves data between two file descriptors without copying between kernel address space and user address space. It transfers up to *len* bytes of data from the file descriptor *fd_in* to the file descriptor *fd_out*, where one of the descriptors must refer to a pipe.

If *fd_in* refers to a pipe, then *off_in* must be NULL. If *fd_in* does not refer to a pipe and *off_in* is NULL, then bytes are read from *fd_in* starting from the current file offset, and the current file offset is adjusted appropriately. If *fd_in* does not refer to a pipe and *off_in* is not NULL, then *off_in* must point to a buffer which specifies the starting offset from which bytes will be read from *fd_in*; in this case, the current file offset of *fd_in* is not changed. Analogous statements apply for *fd_out* and *off_out*.

The *flags* argument is a bit mask that is composed by ORing together zero or more of the following values:

SPLICE_F_MOVE

Attempt to move pages instead of copying. This is only a hint to the kernel: pages may still be copied if the kernel cannot move the pages from the pipe, or if the pipe buffers don't refer to full pages. The initial implementation of this flag was buggy: therefore starting in Linux 2.6.21 it is a no-op (but is still permitted in a **splice()** call); in the future, a correct implementation may be restored.

SPLICE_F_NONBLOCK

Do not block on I/O. This makes the splice pipe operations non-blocking, but **splice()** may nevertheless block because the file descriptors that are spliced to/from may block (unless they have the **O_NONBLOCK** flag set).

SPLICE_F_MORE

More data will be coming in a subsequent splice. This is a helpful hint when the *fd_out* refers to a socket (see also the description of **MSG_MORE** in **send(2)**, and the description of **TCP_CORK** in **tcp(7)**)

SPLICE_F_GIFT

Unused for **splice()**; see **vmsplice(2)**.

RETURN VALUE

Upon successful completion, **splice()** returns the number of bytes spliced to or from the pipe. A return value of 0 means that there was no data to transfer, and it would not make sense to block, because there are no writers connected to the write end of the pipe referred to by *fd_in*.

On error, **splice()** returns -1 and *errno* is set to indicate the error.

ERRORS**EBADF**

One or both file descriptors are not valid, or do not have proper read-write mode.

EINVAL

Target file system doesn't support splicing; target file is opened in append mode; neither of the descriptors refers to a pipe; or offset given for non-seekable device.

ENOMEM

Out of memory.

ESPIPE

Either *off_in* or *off_out* was not NULL, but the corresponding file descriptor refers to a pipe.

VERSIONS

The **splice()** system call first appeared in Linux 2.6.17.

CONFORMING TO

This system call is Linux-specific.

NOTES

The three system calls **splice()**, **vmsplice(2)**, and **tee(2)**, provide userspace programs with full control over an arbitrary kernel buffer, implemented within the kernel using the same type of buffer that is used for a pipe. In overview, these system calls perform the following tasks:

splice()	moves data from the buffer to an arbitrary file descriptor, or vice versa, or from one buffer to another.
tee(2)	"copies" the data from one buffer to another.
vmsplice(2)	"copies" data from user space into the buffer.

Though we talk of copying, actual copies are generally avoided. The kernel does this by implementing a pipe buffer as a set of reference-counted pointers to pages of kernel memory. The kernel creates "copies" of pages in a buffer by creating new pointers (for the output buffer) referring to the pages, and increasing the reference counts for the pages: only pointers are copied, not the pages of the buffer.

EXAMPLE

See **tee(2)**.

SEE ALSO

sendfile(2), **tee(2)**, **vmsplice(2)**, **feature_test_macros(7)**

COLOPHON

This page is part of release 3.22 of the Linux *man-pages* project. A description of the project, and information about reporting bugs, can be found at <http://www.kernel.org/doc/man-pages/>.