

NAME

UTF-8 – an ASCII compatible multi-byte Unicode encoding

DESCRIPTION

The **Unicode 3.0** character set occupies a 16-bit code space. The most obvious Unicode encoding (known as **UCS-2**) consists of a sequence of 16-bit words. Such strings can contain as parts of many 16-bit characters bytes like `'\0'` or `'/'` which have a special meaning in filenames and other C library function arguments. In addition, the majority of Unix tools expects ASCII files and can't read 16-bit words as characters without major modifications. For these reasons, **UCS-2** is not a suitable external encoding of **Unicode** in filenames, text files, environment variables, etc. The **ISO 10646 Universal Character Set (UCS)**, a superset of Unicode, occupies even a 31-bit code space and the obvious **UCS-4** encoding for it (a sequence of 32-bit words) has the same problems.

The **UTF-8** encoding of **Unicode** and **UCS** does not have these problems and is the common way in which **Unicode** is used on Unix-style operating systems.

Properties

The **UTF-8** encoding has the following nice properties:

- * **UCS** characters 0x00000000 to 0x0000007f (the classic **US-ASCII** characters) are encoded simply as bytes 0x00 to 0x7f (ASCII compatibility). This means that files and strings which contain only 7-bit ASCII characters have the same encoding under both **ASCII** and **UTF-8**.
- * All **UCS** characters greater than 0x7f are encoded as a multi-byte sequence consisting only of bytes in the range 0x80 to 0xfd, so no ASCII byte can appear as part of another character and there are no problems with, for example, `'\0'` or `'/'`.
- * The lexicographic sorting order of **UCS-4** strings is preserved.
- * All possible 2³¹ UCS codes can be encoded using **UTF-8**.
- * The bytes 0xfe and 0xff are never used in the **UTF-8** encoding.
- * The first byte of a multi-byte sequence which represents a single non-ASCII **UCS** character is always in the range 0xc0 to 0xfd and indicates how long this multi-byte sequence is. All further bytes in a multi-byte sequence are in the range 0x80 to 0xbf. This allows easy resynchronization and makes the encoding stateless and robust against missing bytes.
- * **UTF-8** encoded **UCS** characters may be up to six bytes long, however the **Unicode** standard specifies no characters above 0x10ffff, so Unicode characters can only be up to four bytes long in **UTF-8**.

Encoding

The following byte sequences are used to represent a character. The sequence to be used depends on the UCS code number of the character:

0x00000000 – 0x0000007F:

0xxxxxx

0x00000080 – 0x000007FF:

110xxxx 10xxxxx

0x00000800 – 0x0000FFFF:

1110xxx 10xxxxx 10xxxxx

0x00010000 – 0x001FFFFF:

11110xxx 10xxxxx 10xxxxx 10xxxxx

0x00200000 – 0x03FFFFFF:

111110xx 10xxxxx 10xxxxx 10xxxxx 10xxxxx

0x04000000 – 0x7FFFFFFF:

1111110x 10xxxxx 10xxxxx 10xxxxx 10xxxxx 10xxxxx

The xxx bit positions are filled with the bits of the character code number in binary representation. Only the shortest possible multi-byte sequence which can represent the code number of the character can be

used.

The **UCS** code values 0xd800–0xdfff (UTF-16 surrogates) as well as 0xfffe and 0xffff (UCS non-characters) should not appear in conforming **UTF-8** streams.

Example

The **Unicode** character 0xa9 = 1010 1001 (the copyright sign) is encoded in UTF-8 as

11000010 10101001 = 0xc2 0xa9

and character 0x2260 = 0010 0010 0110 0000 (the "not equal" symbol) is encoded as:

11100010 10001001 10100000 = 0xe2 0x89 0xa0

Application Notes

Users have to select a **UTF-8** locale, for example with

```
export LANG=en_GB.UTF-8
```

in order to activate the **UTF-8** support in applications.

Application software that has to be aware of the used character encoding should always set the locale with for example

```
setlocale(LC_CTYPE, "")
```

and programmers can then test the expression

```
strcmp(nl_langinfo(CODESET), "UTF-8") == 0
```

to determine whether a **UTF-8** locale has been selected and whether therefore all plaintext standard input and output, terminal communication, plaintext file content, filenames and environment variables are encoded in **UTF-8**.

Programmers accustomed to single-byte encodings such as **US-ASCII** or **ISO 8859** have to be aware that two assumptions made so far are no longer valid in **UTF-8** locales. Firstly, a single byte does not necessarily correspond any more to a single character. Secondly, since modern terminal emulators in **UTF-8** mode also support Chinese, Japanese, and Korean **double-width characters** as well as non-spacing **combining characters**, outputting a single character does not necessarily advance the cursor by one position as it did in **ASCII**. Library functions such as **mbsrtowcs(3)** and **wcswidth(3)** should be used today to count characters and cursor positions.

The official ESC sequence to switch from an **ISO 2022** encoding scheme (as used for instance by VT100 terminals) to **UTF-8** is ESC % G ("x1b%G"). The corresponding return sequence from **UTF-8** to ISO 2022 is ESC % @ ("x1b%@"). Other ISO 2022 sequences (such as for switching the G0 and G1 sets) are not applicable in UTF-8 mode.

It can be hoped that in the foreseeable future, **UTF-8** will replace **ASCII** and **ISO 8859** at all levels as the common character encoding on POSIX systems, leading to a significantly richer environment for handling plain text.

Security

The **Unicode** and **UCS** standards require that producers of **UTF-8** shall use the shortest form possible, for example, producing a two-byte sequence with first byte 0xc0 is non-conforming. **Unicode 3.1** has added the requirement that conforming programs must not accept non-shortest forms in their input. This is for security reasons: if user input is checked for possible security violations, a program might check only for the **ASCII** version of **"/./"** or **";"** or **NUL** and overlook that there are many non-**ASCII** ways to represent these things in a non-shortest **UTF-8** encoding.

Standards

ISO/IEC 10646-1:2000, Unicode 3.1, RFC 2279, Plan 9.

SEE ALSO

nl_langinfo(3), **setlocale(3)**, **charsets(7)**, **unicode(7)**

COLOPHON

This page is part of release 3.22 of the Linux *man-pages* project. A description of the project, and information about reporting bugs, can be found at <http://www.kernel.org/doc/man-pages/>.