

**NAME**

`chown`, `fchown`, `lchown` – change ownership of a file

**SYNOPSIS**

```
#include <unistd.h>
```

```
int chown(const char *path, uid_t owner, gid_t group);
int fchown(int fd, uid_t owner, gid_t group);
int lchown(const char *path, uid_t owner, gid_t group);
```

Feature Test Macro Requirements for glibc (see `feature_test_macros(7)`):

```
fchown(), lchown(): _BSD_SOURCE || _XOPEN_SOURCE >= 500
```

**DESCRIPTION**

These system calls change the owner and group of a file. They differ only in how the file is specified:

- \* **chown()** changes the ownership of the file specified by *path*, which is dereferenced if it is a symbolic link.
- \* **fchown()** changes the ownership of the file referred to by the open file descriptor *fd*.
- \* **lchown()** is like **chown()**, but does not dereference symbolic links.

Only a privileged process (Linux: one with the **CAP\_CHOWN** capability) may change the owner of a file. The owner of a file may change the group of the file to any group of which that owner is a member. A privileged process (Linux: with **CAP\_CHOWN**) may change the group arbitrarily.

If the *owner* or *group* is specified as `-1`, then that ID is not changed.

When the owner or group of an executable file are changed by a non-superuser, the **S\_ISUID** and **S\_ISGID** mode bits are cleared. POSIX does not specify whether this also should happen when root does the **chown()**; the Linux behavior depends on the kernel version. In case of a non-group-executable file (i.e., one for which the **S\_IXGRP** bit is not set) the **S\_ISGID** bit indicates mandatory locking, and is not cleared by a **chown()**.

**RETURN VALUE**

On success, zero is returned. On error, `-1` is returned, and *errno* is set appropriately.

**ERRORS**

Depending on the file system, other errors can be returned. The more general errors for **chown()** are listed below.

**EACCES**

Search permission is denied on a component of the path prefix. (See also **path\_resolution(7)**.)

**EFAULT**

*path* points outside your accessible address space.

**ELOOP**

Too many symbolic links were encountered in resolving *path*.

**ENAMETOOLONG**

*path* is too long.

**ENOENT**

The file does not exist.

**ENOMEM**

Insufficient kernel memory was available.

**ENOTDIR**

A component of the path prefix is not a directory.

**EPERM**

The calling process did not have the required permissions (see above) to change owner and/or group.

**EROFS**

The named file resides on a read-only file system.

The general errors for **fchown()** are listed below:

**EBADF**

The descriptor is not valid.

**EIO**

A low-level I/O error occurred while modifying the inode.

**ENOENT**

See above.

**EPERM**

See above.

**EROFS**

See above.

**CONFORMING TO**

4.4BSD, SVr4, POSIX.1-2001.

The 4.4BSD version can only be used by the superuser (that is, ordinary users cannot give away files).

**NOTES**

When a new file is created (by, for example, **open(2)** or **mkdir(2)**), its owner is made the same as the file system user ID of the creating process. The group of the file depends on a range of factors, including the type of file system, the options used to mount the file system, and whether or not the set-group-ID permission bit is enabled on the parent directory. If the file system supports the *-o grpuid* (or, synonymously *-o bsdgroups*) and *-o nogrpuid* (or, synonymously *-o sysvgroups*) **mount(8)** options, then the rules are as follows:

- \* If the file system is mounted with *-o grpuid*, then the group of a new file is made the same as that of the parent directory.
- \* If the file system is mounted with *-o nogrpuid* and the set-group-ID bit is disabled on the parent directory, then the group of a new file is made the same as the process's file system GID.
- \* If the file system is mounted with *-o nogrpuid* and the set-group-ID bit is enabled on the parent directory, then the group of a new file is made the same as that of the parent directory.

As at Linux 2.6.25, the *-o grpuid* and *-o nogrpuid* mount options are supported by ext2, ext3, ext4, and XFS. File systems that don't support these mount options follow the *-o nogrpuid* rules.

The **chown()** semantics are deliberately violated on NFS file systems which have UID mapping enabled. Additionally, the semantics of all system calls which access the file contents are violated, because **chown()** may cause immediate access revocation on already open files. Client side caching may lead to a delay between the time where ownership have been changed to allow access for a user and the time where the file can actually be accessed by the user on other clients.

In versions of Linux prior to 2.1.81 (and distinct from 2.1.46), **chown()** did not follow symbolic links. Since Linux 2.1.81, **chown()** does follow symbolic links, and there is a new system call **lchown()** that does not follow symbolic links. Since Linux 2.1.86, this new call (that has the same semantics as the old **chown()**) has got the same syscall number, and **chown()** got the newly introduced number.

**EXAMPLE**

The following program changes the ownership of the file named in its second command-line argument to the value specified in its first command-line argument. The new owner can be specified either as a numeric user ID, or as a username (which is converted to a user ID by using **getpwnam**(3) to perform a lookup in the system password file).

```
#include <pwd.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int
main(int argc, char *argv[])
{
    uid_t uid;
    struct passwd *pwd;
    char *endptr;

    if (argc != 3 || argv[1][0] == '\0') {
        fprintf(stderr, "%s <owner> <file>\n", argv[0]);
        exit(EXIT_FAILURE);
    }

    uid = strtol(argv[1], &endptr, 10); /* Allow a numeric string */

    if (*endptr != '\0') { /* Was not pure numeric string */
        pwd = getpwnam(argv[1]); /* Try getting UID for username */
        if (pwd == NULL) {
            perror("getpwnam");
            exit(EXIT_FAILURE);
        }

        uid = pwd->pw_uid;
    }

    if (chown(argv[2], uid, -1) == -1) {
        perror("chown");
        exit(EXIT_FAILURE);
    } /* if */

    exit(EXIT_SUCCESS);
} /* main */
```

**SEE ALSO**

**chmod**(2), **fchownat**(2), **flock**(2), **path\_resolution**(7), **symlink**(7)

**COLOPHON**

This page is part of release 3.22 of the Linux *man-pages* project. A description of the project, and information about reporting bugs, can be found at <http://www.kernel.org/doc/man-pages/>.