

**NAME**

link – make a new name for a file

**SYNOPSIS**

```
#include <unistd.h>
```

```
int link(const char *oldpath, const char *newpath);
```

**DESCRIPTION**

**link()** creates a new link (also known as a hard link) to an existing file.

If *newpath* exists it will *not* be overwritten.

This new name may be used exactly as the old one for any operation; both names refer to the same file (and so have the same permissions and ownership) and it is impossible to tell which name was the "original".

**RETURN VALUE**

On success, zero is returned. On error, `-1` is returned, and *errno* is set appropriately.

**ERRORS****EACCES**

Write access to the directory containing *newpath* is denied, or search permission is denied for one of the directories in the path prefix of *oldpath* or *newpath*. (See also **path\_resolution(7)**.)

**EEXIST**

*newpath* already exists.

**EFAULT**

*oldpath* or *newpath* points outside your accessible address space.

**EIO** An I/O error occurred.

**ELOOP**

Too many symbolic links were encountered in resolving *oldpath* or *newpath*.

**EMLINK**

The file referred to by *oldpath* already has the maximum number of links to it.

**ENAMETOOLONG**

*oldpath* or *newpath* was too long.

**ENOENT**

A directory component in *oldpath* or *newpath* does not exist or is a dangling symbolic link.

**ENOMEM**

Insufficient kernel memory was available.

**ENOSPC**

The device containing the file has no room for the new directory entry.

**ENOTDIR**

A component used as a directory in *oldpath* or *newpath* is not, in fact, a directory.

**EPERM**

*oldpath* is a directory.

**EPERM**

The file system containing *oldpath* and *newpath* does not support the creation of hard links.

**EROFS**

The file is on a read-only file system.

**EXDEV**

*oldpath* and *newpath* are not on the same mounted file system. (Linux permits a file system to be mounted at multiple points, but **link()** does not work across different mount points, even if the

same file system is mounted on both.)

## CONFORMING TO

SVr4, 4.3BSD, POSIX.1-2001 (but see NOTES).

## NOTES

Hard links, as created by **link()**, cannot span file systems. Use **symlink(2)** if this is required.

POSIX.1-2001 says that **link()** should dereference *oldpath* if it is a symbolic link. However, since kernel 2.0, Linux does not do so: if *oldpath* is a symbolic link, then *newpath* is created as a (hard) link to the same symbolic link file (i.e., *newpath* becomes a symbolic link to the same file that *oldpath* refers to). Some other implementations behave in the same manner as Linux. POSIX.1-2008 changes the specification of **link()**, making it implementation-dependent whether or not *oldpath* is dereferenced if it is a symbolic link. For precise control over the treatment of symbolic links when creating a link, see **linkat(2)**.

## BUGS

On NFS file systems, the return code may be wrong in case the NFS server performs the link creation and dies before it can say so. Use **stat(2)** to find out if the link got created.

## SEE ALSO

**ln(1)**, **linkat(2)**, **open(2)**, **rename(2)**, **stat(2)**, **symlink(2)**, **unlink(2)**, **path\_resolution(7)**, **symlink(7)**

## COLOPHON

This page is part of release 3.22 of the Linux *man-pages* project. A description of the project, and information about reporting bugs, can be found at <http://www.kernel.org/doc/man-pages/>.