

**NAME**

semctl – semaphore control operations

**SYNOPSIS**

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
int semctl(int semid, int semnum, int cmd, ...);
```

**DESCRIPTION**

**semctl()** performs the control operation specified by *cmd* on the semaphore set identified by *semid*, or on the *semnum*-th semaphore of that set. (The semaphores in a set are numbered starting at 0.)

This function has three or four arguments, depending on *cmd*. When there are four, the fourth has the type *union semun*. The *calling program* must define this union as follows:

```
union semun {
    int      val; /* Value for SETVAL */
    struct semid_ds *buf; /* Buffer for IPC_STAT, IPC_SET */
    unsigned short *array; /* Array for GETALL, SETALL */
    struct seminfo *__buf; /* Buffer for IPC_INFO
                           (Linux-specific) */
};
```

The *semid\_ds* data structure is defined in *<sys/sem.h>* as follows:

```
struct semid_ds {
    struct ipc_perm sem_perm; /* Ownership and permissions */
    time_t      sem_otime; /* Last semop time */
    time_t      sem_ctime; /* Last change time */
    unsigned short sem_nsems; /* No. of semaphores in set */
};
```

The *ipc\_perm* structure is defined in *<sys/ipc.h>* as follows (the highlighted fields are settable using **IPC\_SET**):

```
struct ipc_perm {
    key_t      __key; /* Key supplied to semget(2) */
    uid_t      uid; /* Effective UID of owner */
    gid_t      gid; /* Effective GID of owner */
    uid_t      cuid; /* Effective UID of creator */
    gid_t      cgid; /* Effective GID of creator */
    unsigned short mode; /* Permissions */
    unsigned short __seq; /* Sequence number */
};
```

Valid values for *cmd* are:

**IPC\_STAT**

Copy information from the kernel data structure associated with *semid* into the *semid\_ds* structure pointed to by *arg.buf*. The argument *semnum* is ignored. The calling process must have read permission on the semaphore set.

**IPC\_SET** Write the values of some members of the *semid\_ds* structure pointed to by *arg.buf* to the kernel data structure associated with this semaphore set, updating also its *sem\_ctime* member. The following members of the structure are updated: *sem\_perm.uid*, *sem\_perm.gid*, and (the least significant 9 bits of) *sem\_perm.mode*. The effective UID of the calling process must match the owner (*sem\_perm.uid*) or creator (*sem\_perm.cuid*) of the semaphore set, or the caller must be privileged. The argument *semnum* is ignored.

**IPC\_RMID**

Immediately remove the semaphore set, awakening all processes blocked in **semop**(2) calls on the set (with an error return and *errno* set to **EIDRM**). The effective user ID of the calling process must match the creator or owner of the semaphore set, or the caller must be privileged. The argument *semnum* is ignored.

**IPC\_INFO** (Linux-specific)

Returns information about system-wide semaphore limits and parameters in the structure pointed to by *arg.\_\_buf*. This structure is of type *seminfo*, defined in *<sys/sem.h>* if the **\_GNU\_SOURCE** feature test macro is defined:

```
struct seminfo {
    int semmap; /* Number of entries in semaphore
                map; unused within kernel */
    int semmni; /* Maximum number of semaphore sets */
    int semmns; /* Maximum number of semaphores in all
                semaphore sets */
    int semmnu; /* System-wide maximum number of undo
                structures; unused within kernel */
    int semmsl; /* Maximum number of semaphores in a
                set */
    int semopm; /* Maximum number of operations for
                semop(2) */
    int semume; /* Maximum number of undo entries per
                process; unused within kernel */
    int semusz; /* Size of struct sem_undo */
    int semvmx; /* Maximum semaphore value */
    int semaem; /* Max. value that can be recorded for
                semaphore adjustment (SEM_UNDO) */
};
```

The *semmsl*, *semmns*, *semopm*, and *semmni* settings can be changed via */proc/sys/kernel/sem*; see **proc**(5) for details.

**SEM\_INFO** (Linux-specific)

Returns a *seminfo* structure containing the same information as for **IPC\_INFO**, except that the following fields are returned with information about system resources consumed by semaphores: the *semusz* field returns the number of semaphore sets that currently exist on the system; and the *semaem* field returns the total number of semaphores in all semaphore sets on the system.

**SEM\_STAT** (Linux-specific)

Returns a *semid\_ds* structure as for **IPC\_STAT**. However, the *semid* argument is not a semaphore identifier, but instead an index into the kernel's internal array that maintains information about all semaphore sets on the system.

**GETALL** Return **semval** (i.e., the current value) for all semaphores of the set into *arg.array*. The argument *semnum* is ignored. The calling process must have read permission on the semaphore set.

**GETNCNT**

The system call returns the value of **semncnt** (i.e., the number of processes waiting for the value of this semaphore to increase) for the *semnum*-th semaphore of the set (i.e., the number of processes waiting for an increase of **semval** for the *semnum*-th semaphore of the set). The calling process must have read permission on the semaphore set.

**GETPID** The system call returns the value of **sempid** for the *semnum*-th semaphore of the set (i.e., the PID of the process that executed the last **semop**(2) call for the *semnum*-th semaphore of the set). The calling process must have read permission on the semaphore set.

**GETVAL** The system call returns the value of **semval** for the *semnum*-th semaphore of the set. The calling process must have read permission on the semaphore set.

**GETZCNT**

The system call returns the value of **semzcnt** (i.e., the number of processes waiting for the value of this semaphore to become zero) for the *semnum*-th semaphore of the set (i.e., the number of processes waiting for **semval** of the *semnum*-th semaphore of the set to become 0). The calling process must have read permission on the semaphore set.

**SETALL** Set **semval** for all semaphores of the set using *arg.array*, updating also the *sem\_ctime* member of the *semid\_ds* structure associated with the set. Undo entries (see **semop**(2)) are cleared for altered semaphores in all processes. If the changes to semaphore values would permit blocked **semop**(2) calls in other processes to proceed, then those processes are woken up. The argument *semnum* is ignored. The calling process must have alter (write) permission on the semaphore set.

**SETVAL** Set the value of **semval** to *arg.val* for the *semnum*-th semaphore of the set, updating also the *sem\_ctime* member of the *semid\_ds* structure associated with the set. Undo entries are cleared for altered semaphores in all processes. If the changes to semaphore values would permit blocked **semop**(2) calls in other processes to proceed, then those processes are woken up. The calling process must have alter permission on the semaphore set.

## RETURN VALUE

On failure **semctl**() returns -1 with *errno* indicating the error.

Otherwise the system call returns a non-negative value depending on *cmd* as follows:

**GETNCNT** the value of **semncnt**.

**GETPID** the value of **sempid**.

**GETVAL** the value of **semval**.

**GETZCNT** the value of **semzcnt**.

**IPC\_INFO** the index of the highest used entry in the kernel's internal array recording information about all semaphore sets. (This information can be used with repeated **SEM\_STAT** operations to obtain information about all semaphore sets on the system.)

**SEM\_INFO** As for **IPC\_INFO**.

**SEM\_STAT** the identifier of the semaphore set whose index was given in *semid*.

All other *cmd* values return 0 on success.

## ERRORS

On failure, *errno* will be set to one of the following:

**EACCES**

The argument *cmd* has one of the values **GETALL**, **GETPID**, **GETVAL**, **GETNCNT**, **GETZCNT**, **IPC\_STAT**, **SEM\_STAT**, **SETALL**, or **SETVAL** and the calling process does not have the required permissions on the semaphore set and does not have the **CAP\_IPC\_OWNER** capability.

**EFAULT**

The address pointed to by *arg.buf* or *arg.array* isn't accessible.

**EIDRM**

The semaphore set was removed.

**EINVAL**

Invalid value for *cmd* or *semid*. Or: for a **SEM\_STAT** operation, the index value specified in *semid* referred to an array slot that is currently unused.

**EPERM**

The argument *cmd* has the value **IPC\_SET** or **IPC\_RMID** but the effective user ID of the calling process is not the creator (as found in *sem\_perm.cuid*) or the owner (as found in *sem\_perm.uid*) of the semaphore set, and the process does not have the **CAP\_SYS\_ADMIN** capability.

**ERANGE**

The argument *cmd* has the value **SETALL** or **SETVAL** and the value to which **semval** is to be set (for some semaphore of the set) is less than 0 or greater than the implementation limit **SEMVMX**.

**CONFORMING TO**

SVr4, POSIX.1-2001.

**NOTES**

The **IPC\_INFO**, **SEM\_STAT** and **SEM\_INFO** operations are used by the **ipcs**(8) program to provide information on allocated resources. In the future these may be modified or moved to a /proc file system interface.

Various fields in a *struct semid\_ds* were typed as *short* under Linux 2.2 and have become *long* under Linux 2.4. To take advantage of this, a recompilation under glibc-2.1.91 or later should suffice. (The kernel distinguishes old and new calls by an **IPC\_64** flag in *cmd*.)

In some earlier versions of glibc, the *semun* union was defined in *<sys/sem.h>*, but POSIX.1-2001 requires that the caller define this union. On versions of glibc where this union is *not* defined, the macro **\_SEM\_SEMUN\_UNDEFINED** is defined in *<sys/sem.h>*.

The following system limit on semaphore sets affects a **semctl**() call:

**SEMVMX**

Maximum value for **semval**: implementation dependent (32767).

For greater portability it is best to always call **semctl**() with four arguments.

**SEE ALSO**

**ipc**(2), **semget**(2), **semop**(2), **capabilities**(7), **sem\_overview**(7), **svipc**(7)

**COLOPHON**

This page is part of release 3.22 of the Linux *man-pages* project. A description of the project, and information about reporting bugs, can be found at <http://www.kernel.org/doc/man-pages/>.