

NAME

statfs, fstatfs – get filesystem statistics

SYNOPSIS

```
#include <sys/vfs.h> /* or <sys/statfs.h> */
```

```
int statfs(const char *path, struct statfs *buf);
```

```
int fstatfs(int fd, struct statfs *buf);
```

DESCRIPTION

The **statfs()** system call returns information about a mounted filesystem. *path* is the pathname of any file within the mounted filesystem. *buf* is a pointer to a *statfs* structure defined approximately as follows:

```
struct statfs {
    __fsword_t f_type; /* Type of filesystem (see below) */
    __fsword_t f_bsize; /* Optimal transfer block size */
    fsblkcnt_t f_blocks; /* Total data blocks in filesystem */
    fsblkcnt_t f_bfree; /* Free blocks in filesystem */
    fsblkcnt_t f_bavail; /* Free blocks available to
                          unprivileged user */
    fsfilcnt_t f_files; /* Total file nodes in filesystem */
    fsfilcnt_t f_ffree; /* Free file nodes in filesystem */
    fsid_t f_fsid; /* Filesystem ID */
    __fsword_t f_namelen; /* Maximum length of filenames */
    __fsword_t f_frsize; /* Fragment size (since Linux 2.6) */
    __fsword_t f_flags; /* Mount flags of filesystem
                        (since Linux 2.6.36) */
    __fsword_t f_spare[xxx];
    /* Padding bytes reserved for future use */
};
```

Filesystem types:

```
ADFS_SUPER_MAGIC    0xadf5
AFFS_SUPER_MAGIC    0xADFF
BDEVFS_MAGIC        0x62646576
BEFS_SUPER_MAGIC    0x42465331
BFS_MAGIC           0x1BADFACE
BINFMTFS_MAGIC      0x42494e4d
BTRFS_SUPER_MAGIC   0x9123683E
CGROUP_SUPER_MAGIC  0x27e0eb
CIFS_MAGIC_NUMBER    0xFF534D42
CODA_SUPER_MAGIC    0x73757245
COH_SUPER_MAGIC     0x012FF7B7
CRAMFS_MAGIC        0x28cd3d45
DEBUGFS_MAGIC       0x64626720
DEVFS_SUPER_MAGIC   0x1373
DEVPTS_SUPER_MAGIC  0x1cd1
EFIVARFS_MAGIC      0xde5e81e4
EFS_SUPER_MAGIC     0x00414A53
EXT_SUPER_MAGIC     0x137D
EXT2_OLD_SUPER_MAGIC 0xEF51
EXT2_SUPER_MAGIC    0xEF53
EXT3_SUPER_MAGIC    0xEF53
EXT4_SUPER_MAGIC    0xEF53
FUSE_SUPER_MAGIC     0x6573546
```

```

FUTEXFS_SUPER_MAGIC 0xBAD1DEA
HFS_SUPER_MAGIC     0x4244
HOSTFS_SUPER_MAGIC   0x00c0ffee
HPFS_SUPER_MAGIC     0xF995E849
HUGETLBFS_MAGIC      0x958458f6
ISOFS_SUPER_MAGIC    0x9660
JFFS2_SUPER_MAGIC    0x72b6
JFS_SUPER_MAGIC       0x3153464a
MINIX_SUPER_MAGIC     0x137F /* orig. minix */
MINIX_SUPER_MAGIC2    0x138F /* 30 char minix */
MINIX2_SUPER_MAGIC    0x2468 /* minix V2 */
MINIX2_SUPER_MAGIC2   0x2478 /* minix V2, 30 char names */
MINIX3_SUPER_MAGIC    0x4d5a /* minix V3 fs, 60 char names */
MQQUEUE_MAGIC         0x19800202
MSDOS_SUPER_MAGIC     0x4d44
NCP_SUPER_MAGIC       0x564c
NFS_SUPER_MAGIC       0x6969
NILFS_SUPER_MAGIC     0x3434
NTFS_SB_MAGIC         0x5346544e
OCFS2_SUPER_MAGIC     0x7461636f
OPENPROM_SUPER_MAGIC 0x9fa1
PIPEFS_MAGIC          0x50495045
PROC_SUPER_MAGIC      0x9fa0
PSTOREFS_MAGIC        0x6165676C
QNX4_SUPER_MAGIC      0x002f
QNX6_SUPER_MAGIC      0x68191122
RAMFS_MAGIC           0x858458f6
REISERFS_SUPER_MAGIC 0x52654973
ROMFS_MAGIC           0x7275
SELINUX_MAGIC         0xf97cff8c
SMACK_MAGIC           0x43415d53
SMB_SUPER_MAGIC       0x517B
SOCKFS_MAGIC          0x534F434B
SQUASHFS_MAGIC        0x73717368
SYSFS_MAGIC           0x62656572
SYSV2_SUPER_MAGIC     0x012FF7B6
SYSV4_SUPER_MAGIC     0x012FF7B5
TMPFS_MAGIC           0x01021994
UDF_SUPER_MAGIC       0x15013346
UFS_MAGIC             0x00011954
USBDEVICE_SUPER_MAGIC 0x9fa2
V9FS_MAGIC            0x01021997
VXFS_SUPER_MAGIC      0xa501FCF5
XENFS_SUPER_MAGIC     0xabba1974
XENIX_SUPER_MAGIC     0x012FF7B4
XFS_SUPER_MAGIC       0x58465342
_XIAFS_SUPER_MAGIC    0x012FD16D

```

Most of these MAGIC constants are defined in `/usr/include/linux/magic.h`, and some are hardcoded in kernel sources.

The `f_flags` is a bit mask indicating mount options for the file system. It contains zero or more of the following bits:

ST_MANDLOCK

Mandatory locking is permitted on the filesystem (see **fcntl(2)**).

ST_NOATIME

Do not update access times; see **mount(2)**.

ST_NODEV

Disallow access to device special files on this filesystem.

ST_NODIRATIME

Do not update directory access times; see **mount(2)**.

ST_NOEXEC

Execution of programs is disallowed on this filesystem.

ST_NOSUID

The set-user-ID and set-group-ID bits are ignored by **exec(3)** for executable files on this filesystem

ST_RDONLY

This filesystem is mounted read-only.

ST_RELATIME

Update atime relative to mtime/ctime; see **mount(2)**.

ST_SYNCHRONOUS

Writes are synched to the filesystem immediately (see the description of **O_SYNC** in **open(2)**).

Nobody knows what *f_fsid* is supposed to contain (but see below).

Fields that are undefined for a particular filesystem are set to 0.

fstatfs() returns the same information about an open file referenced by descriptor *fd*.

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and *errno* is set appropriately.

ERRORS**EACCES**

(**statfs()**) Search permission is denied for a component of the path prefix of *path*. (See also **path_resolution(7)**.)

EBADF

(**fstatfs()**) *fd* is not a valid open file descriptor.

EFAULT

buf or *path* points to an invalid address.

EINTR

This call was interrupted by a signal.

EIO

An I/O error occurred while reading from the filesystem.

ELOOP

(**statfs()**) Too many symbolic links were encountered in translating *path*.

ENAMETOOLONG

(**statfs()**) *path* is too long.

ENOENT

(**statfs()**) The file referred to by *path* does not exist.

ENOMEM

Insufficient kernel memory was available.

ENOSYS

The filesystem does not support this call.

ENOTDIR

(**statfs()**) A component of the path prefix of *path* is not a directory.

EOVERFLOW

Some values were too large to be represented in the returned struct.

CONFORMING TO

Linux-specific. The Linux **statfs()** was inspired by the 4.4BSD one (but they do not use the same structure).

NOTES

The `__fsword_t` type used for various fields in the *statfs* structure definition is a glibc internal type, not intended for public use. This leaves the programmer in a bit of a conundrum when trying to copy or compare these fields to local variables in a program. Using *unsigned int* for such variables suffices on most systems.

The original Linux **statfs()** and **fstatfs()** system calls were not designed with extremely large file sizes in mind. Subsequently, Linux 2.6 added new **statfs64()** and **fstatfs64()** system calls that employ a new structure, *statfs64*. The new structure contains the same fields as the original *statfs* structure, but the sizes of various fields are increased, to accommodate large file sizes. The glibc **statfs()** and **fstatfs()** wrapper functions transparently deal with the kernel differences.

Some systems have only `<sys/vfs.h>`, other systems also have `<sys/statfs.h>`, where the former includes the latter. So it seems including the former is the best choice.

LSB has deprecated the library calls **statfs()** and **fstatfs()** and tells us to use **statvfs(2)** and **fstatvfs(2)** instead.

The *f_fsid* field

Solaris, Irix and POSIX have a system call **statvfs(2)** that returns a *struct statvfs* (defined in `<sys/statvfs.h>`) containing an *unsigned long f_fsid*. Linux, SunOS, HP-UX, 4.4BSD have a system call **statfs()** that returns a *struct statfs* (defined in `<sys/vfs.h>`) containing a *fsid_t f_fsid*, where *fsid_t* is defined as *struct { int val[2]; }*. The same holds for FreeBSD, except that it uses the include file `<sys/mount.h>`.

The general idea is that *f_fsid* contains some random stuff such that the pair (*f_fsid*, *ino*) uniquely determines a file. Some operating systems use (a variation on) the device number, or the device number combined with the filesystem type. Several operating systems restrict giving out the *f_fsid* field to the superuser only (and zero it for unprivileged users), because this field is used in the filehandle of the filesystem when NFS-exported, and giving it out is a security concern.

Under some operating systems, the *fsid* can be used as the second argument to the **sysfs(2)** system call.

BUGS

From Linux 2.6.38 up to and including Linux 3.1, **fstatfs()** failed with the error **ENOSYS** for file descriptors created by **pipe(2)**.

SEE ALSO

stat(2), **statvfs(2)**, **path_resolution(7)**