

NAME

swapon, swapoff – start/stop swapping to file/device

SYNOPSIS

```
#include <unistd.h>
#include <asm/page.h> /* to find PAGE_SIZE */
#include <sys/swap.h>
```

```
int swapon(const char *path, int swapflags);
int swapoff(const char *path);
```

DESCRIPTION

swapon() sets the swap area to the file or block device specified by *path*. **swapoff()** stops swapping to the file or block device specified by *path*.

swapon() takes a *swapflags* argument. If *swapflags* has the **SWAP_FLAG_PREFER** bit turned on, the new swap area will have a higher priority than default. The priority is encoded within *swapflags* as:

$(prio \ll SWAP_FLAG_PRIO_SHIFT) \& SWAP_FLAG_PRIO_MASK$

These functions may only be used by a privileged process (one having the **CAP_SYS_ADMIN** capability).

Priority

Each swap area has a priority, either high or low. The default priority is low. Within the low-priority areas, newer areas are even lower priority than older areas.

All priorities set with *swapflags* are high-priority, higher than default. They may have any non-negative value chosen by the caller. Higher numbers mean higher priority.

Swap pages are allocated from areas in priority order, highest priority first. For areas with different priorities, a higher-priority area is exhausted before using a lower-priority area. If two or more areas have the same priority, and it is the highest priority available, pages are allocated on a round-robin basis between them.

As of Linux 1.3.6, the kernel usually follows these rules, but there are exceptions.

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and *errno* is set appropriately.

ERRORS**EBUSY**

(for **swapon()**) The specified *path* is already being used as a swap area.

EINVAL

The file *path* exists, but refers neither to a regular file nor to a block device; or, for **swapon()**, the indicated path does not contain a valid swap signature or resides on an in-memory file system like tmpfs; or, for **swapoff()**, *path* is not currently a swap area.

ENFILE

The system limit on the total number of open files has been reached.

ENOENT

The file *path* does not exist.

ENOMEM

The system has insufficient memory to start swapping.

EPERM

The caller does not have the **CAP_SYS_ADMIN** capability. Alternatively, the maximum number of swap files are already in use; see NOTES below.

CONFORMING TO

These functions are Linux-specific and should not be used in programs intended to be portable. The second *swapflags* argument was introduced in Linux 1.3.2.

NOTES

The partition or path must be prepared with **mkswap**(8).

There is an upper limit on the number of swap files that may be used, defined by the kernel constant **MAX_SWAPFILES**. Before kernel 2.4.10, **MAX_SWAPFILES** has the value 8; since kernel 2.4.10, it has the value 32. Since kernel 2.6.18, the limit is decreased by 2 if the kernel is built with the **CONFIG_MIGRATION** option (which reserves two swap table entries for the page migration features of **mbind**(2) and **migrate_pages**(2)). Since kernel 2.6.32, the limit is decreased by 1 if the kernel is built with the **CONFIG_MEMORY_FAILURE** option.

SEE ALSO

mkswap(8), **swapoff**(8), **swapon**(8)

COLOPHON

This page is part of release 3.22 of the Linux *man-pages* project. A description of the project, and information about reporting bugs, can be found at <http://www.kernel.org/doc/man-pages/>.